TET Network

A Fluid Peer-to-Peer Computational and Energy Resource Protocol

(The AI-Native Sovereign L1)

Date: April 28, 2026 (Genesis Draft: March 28, 2026)

Author & Architecture: Steve > Founder-Architect > TET Network Project

Contact: yizhenxianshi@gmail.com

1. Foreword

To many, the legacy infrastructure of centralized cloud computing is an unchangeable reality; to me, it is a historical anomaly. Unburdened by the dogmas of the past, I have spent the past few years observing the fundamental contradiction of our digital age: the monopoly of human intelligence by capital. In Singapore, I witnessed the physical reality of this monopoly—massive, opaque data centers dictating the price of progress. In Switzerland, I discovered the antidote—the resilience of decentralized trust and the absolute clarity of mathematical law. My previous architecture established a single, undeniable truth: Compute is Energy. However, traditional blockchains suffer from "Consensus Decay." Satoshi Nakamoto's design forces every participant to carry the same heavy rock, limiting participation to those with massive capital. TET Network is the evolution. It is built on a new realization: The chain must adapt to the machine, not the other way around. By introducing Context-Aware Adaptive Consensus (CAAC), TET Network identifies the environment of the node—whether it is a massive GPU farm or a smartphone—and dynamically changes the rules of mining. We are creating a "Fluid Chain" that turns all 8 billion devices on Earth into a single, indestructible, and quantum-resistant compute entity. I serve as the Founder-Architect. I have provided the definitive map. This is not a job offer. It is a call to "Builders"—engineers who refuse to work for the giants of the past and are ready to redefine intelligence as a common resource for all humanity.

Abstract & Executive Summary: TET at a Glance

An autonomous decentralized network where participants directly share computational power and electricity without central administrators enables the free distribution of resources. Currently, the infrastructure supporting artificial intelligence relies on absolute trust in a very small number of platform companies, creating fatal vulnerabilities: opaque censorship, high intermediary costs, and the risk of system-wide outages.

TET Network is a Post-Quantum, AI-Native Sovereign Layer 1 designed to democratize intelligence through a fundamentally new architecture:

- **Context-Aware Adaptive Consensus (Fluid Chain):** TET dynamically adjusts its validation rules based on the hardware profile of the node. High-end GPU farms execute *Proof of Compute* for high rewards, while edge devices execute lightweight *Proof of Relay* for smaller rewards.
- **Universal Miner Inclusion:** By adapting to the environment, we eliminate the hardware barrier, moving from a 0.01% elite network to a 100% human-inclusive network.
- **Thermodynamic Efficiency:** Meaningless PoW puzzles are eliminated. Every watt of power consumed by a TET miner is used to process a real AI inference or a cryptographic proof.
- **Indestructible Resilience (The Cockroach Doctrine):** If centralized data centers go offline, millions of lightweight mobile nodes automatically take over state propagation. The network cannot be physically killed.
- **Post-Quantum Security:** TET implements ML-DSA (Dilithium) as its primary signature scheme, securing sovereign assets against future quantum threats.

2. Introduction: The Structural Failure of Homogeneous Consensus

The fundamental flaw in current Layer 1 architectures is the assumption of homogeneous node participation. In traditional Proof of Work (PoW), all participants compete in the exact same cryptographic hashing algorithm, logically centralizing processing power into the hands of specialized hardware (ASIC) operators. Proof of Stake (PoS) systems merely replace hardware monopolies with capital monopolies.

In both paradigms, the immense energy and resources consumed do not contribute to external utility or real-world computational progress. TET Network addresses this "Consensus Decay" by decoupling the consensus mechanism from a rigid hardware requirement, transforming external computational demand (AI inference) into the protocol's underlying security model.

3. Core Innovation: Context-Aware Adaptive Consensus (CAAC)

TET Network introduces CAAC, a protocol-level mechanism that dynamically evaluates the physical capabilities and network context of a connecting node, assigning it an optimal validation role. The protocol functions under the principle that the chain must adapt to the physical hardware, rather than forcing hardware to adapt to a single chain rule.

3.1 High-Performance Nodes (Proof of Compute)

Nodes equipped with substantial processing power (e.g., data center GPU clusters) are categorized to execute *Proof of Compute (PoC)*. These nodes process complex, real-world AI inferences (such as LLM matrix operations). The network verifies their output through optimistic execution and Zero-Knowledge (ZK) fraud proofs. Compensation (TET) is issued in direct proportion to the verified thermodynamic energy provided.

3.2 Edge Nodes (Proof of Relay & PQC Validation)

Nodes with constrained computational capacity (e.g., consumer smartphones, IoT devices) are categorized to execute *Proof of Relay (PoR)*. They do not run heavy AI compute. Instead, they sustain the network by routing data, maintaining state telemetry, and executing lightweight verification of Post-Quantum Cryptographic (PQC) signatures.

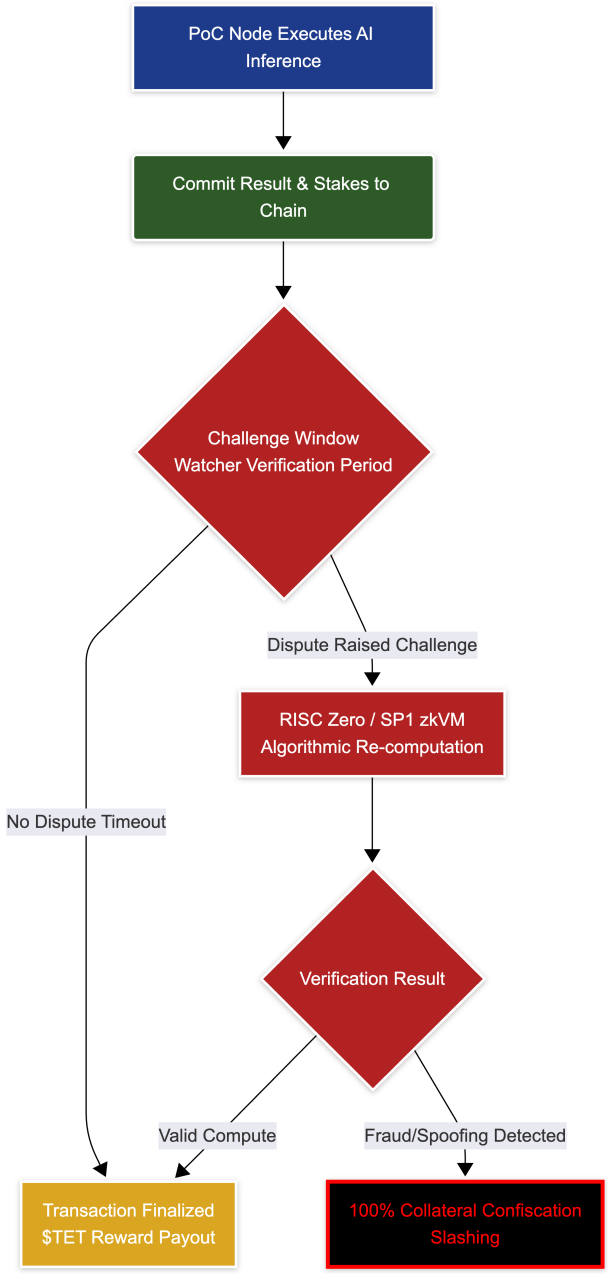

4. The Fluid Architecture and Sovereign Runtime

The defining innovation of TET Network is the Environment-Adaptive Chain (Fluid Chain). The network eliminates abstract benchmarking scores, utilizing **TET** as the universal unit of both value and physical resource. 1 TET represents a deterministic quantity of verified AI computation or network sustainment work.

4.1 The Sovereign Runtime

Ethereum utilizes "Gas" to limit compute steps and prevent infinite loops. TET Network utilizes TET directly to price thermodynamic execution. When a smart contract requests an AI inference, the Sovereign Runtime shifts to an Optimistic Verification model. The task is routed to a PoC node, executed off-main-thread, and the result is returned to the state with a cryptographic commitment. Disputes invoke the ZK-Court (utilizing RISC Zero or SP1 zkVMs) to algorithmically slash malicious actors.

4.2 Mathematical Model of Thermodynamic Work

The fundamental value of 1 TET is not arbitrary; it is a cryptographic proof of physical work. The total expected reward R for a node over a time interval t is calculated based on its verified computational throughput $(C_{flops}(t))$ and thermodynamic energy expenditure $(E_{joules}(t))$, adjusted by the network's dynamic difficulty factor Γ:

$$´R = \int_0^t \left(\frac{C_{flops}(t)}{E_{joules}(t)} \cdot \Gamma \right) dt$$

This equation establishes the ultimate sovereign peg: intelligence generation cryptographically bound to physical electricity.

5. Anatomy of a Fluid Transaction

In the TET Network, a transaction can represent either a transfer of value or a request for physical AI computation. Therefore, the transaction structure must inform the Context-Aware Adaptive Consensus (CAAC) on how to route the data.

A standard TET transaction contains a **Workload Flag**: a binary indicator defining the transaction type. `0` represents a Standard Value Transfer, and `1` represents an AI Inference/Compute Request. If the flag is `1`, the CAAC algorithm exclusively routes it to PoC High-Performance Nodes. If the transaction is type `0`, it is lightly validated by Edge Nodes via PoR.

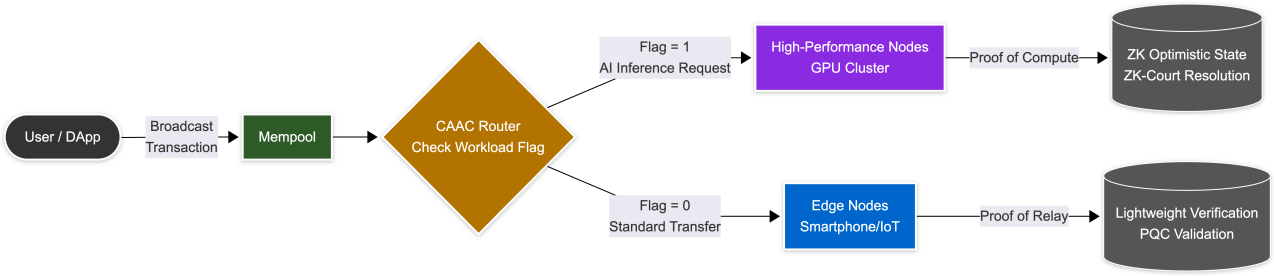

6. Edge Clustering and Semantic Routing

Geographic dispersion of computational resources provides superior censorship resistance but introduces physical network latency. TET Network resolves this through dynamic edge clustering and Semantic Routing.

Large AI model weights (e.g., 50GB tensor files) are not distributed randomly. Instead, nodes specialize and cache specific model families based on their hardware. When a user requests an inference, the protocol dynamically routes the task to the nearest logical node cluster that already possesses the pre-loaded weights (Weight Locality). This eliminates massive data transfer overhead, ensuring decentralized inference speeds match or exceed centralized cloud providers.

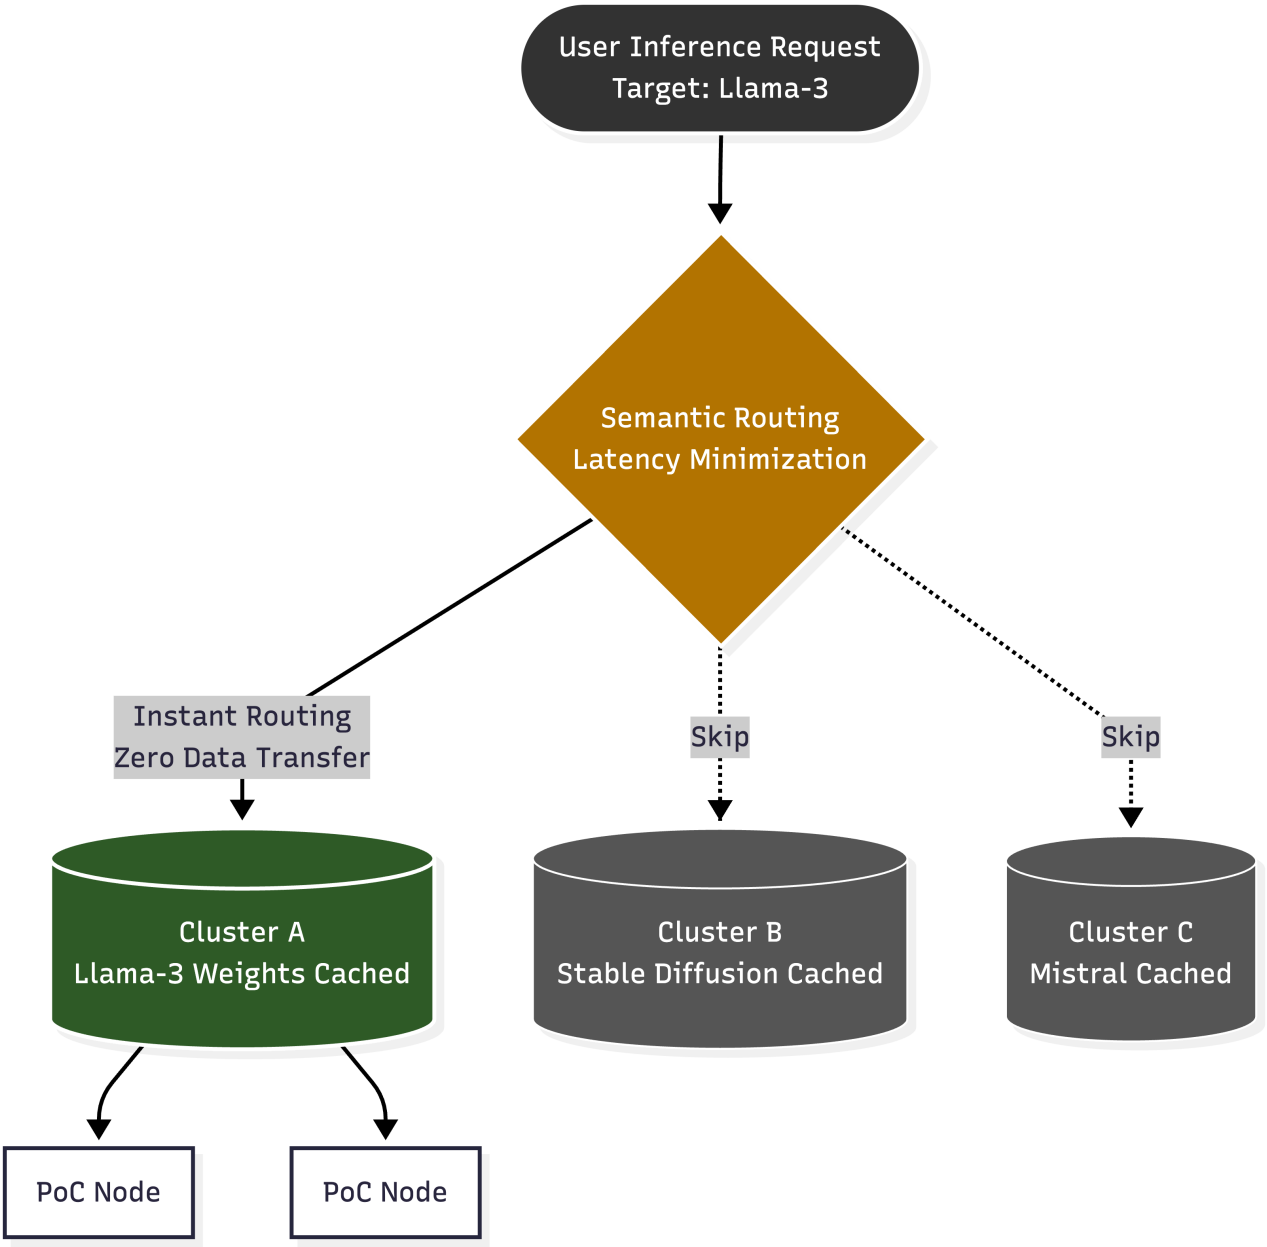

7. State Verification for Edge Devices

The fundamental premise of the "8 Billion Device Plan" requires that smartphones and IoT devices can participate securely without storing the entire ledger history.

TET Network utilizes an advanced multi-level Merkle tree structure to compress verified data. While high-performance PoC nodes store the complete state, Edge Nodes (PoR) operate as Light Clients. They download only block headers and verify the proofs associated with them. When an edge device needs to query a specific contract or balance, it retrieves only the relevant Merkle branch, mathematically proving data validity without needing the full multi-terabyte blockchain.

8. Thermodynamic Efficiency and Network Resilience

TET Network implements a Zero-Waste computational model. Every unit of energy expended by a PoC node is directly allocated to resolving external AI tasks (e.g., neural network inference).

Furthermore, while high-performance nodes handle compute-intensive tasks, the continuous liveness and state propagation of the network are maintained by millions of globally distributed edge nodes. Even if major data centers or GPU clusters unexpectedly go offline, the surrounding edge nodes autonomously sustain network telemetry and routing protocols. This localized redundancy guarantees absolute ledger survivability regardless of physical or geopolitical constraints.

9. Post-Quantum Security

As quantum computing advances, cryptographic standards currently protecting 99% of global blockchain assets (e.g., ECDSA) will become vulnerable to Shor's algorithm. TET Network mitigates this existential risk by implementing Post-Quantum Cryptography (PQC) at the base layer.

The protocol adopts ML-DSA (Dilithium) as its primary signature scheme. This sovereign security model ensures that private keys and on-chain assets within the TET ecosystem are mathematically protected against both classical computers and future quantum adversaries from the genesis state.

10. Tokenomics: Genesis Allocation & Scarcity

To ensure absolute transparency and economic reality, the genesis allocation is hard-coded into the base layer. The maximum total supply is strictly fixed at **10,000,000,000 $TET** (10 Billion).

- **25% - Founder & Core Contributors:** Allocated to the Founder-Architect and core development team. To guarantee long-term commitment to the network's success, these funds are subject to a strict multi-year vesting schedule.
- **75% - Autonomous Network Allocation:** This absolute majority of the supply is intentionally not allocated to any central entity. Instead, it is locked within the protocol's smart contracts to run the network autonomously:
    - **50% (Resource Mining):** Programmatically emitted over decades to reward nodes performing Proof of Compute (PoC) and Proof of Relay (PoR).
    - **25% (Ecosystem & Treasury):** Reserved in a decentralized treasury to fund future AI infrastructure grants and network expansion.
- **Deflationary Burn Mechanism:** To suppress inflation and enforce scarcity, 50% of all transaction fees (including the cost of AI inference tasks) are permanently burned.

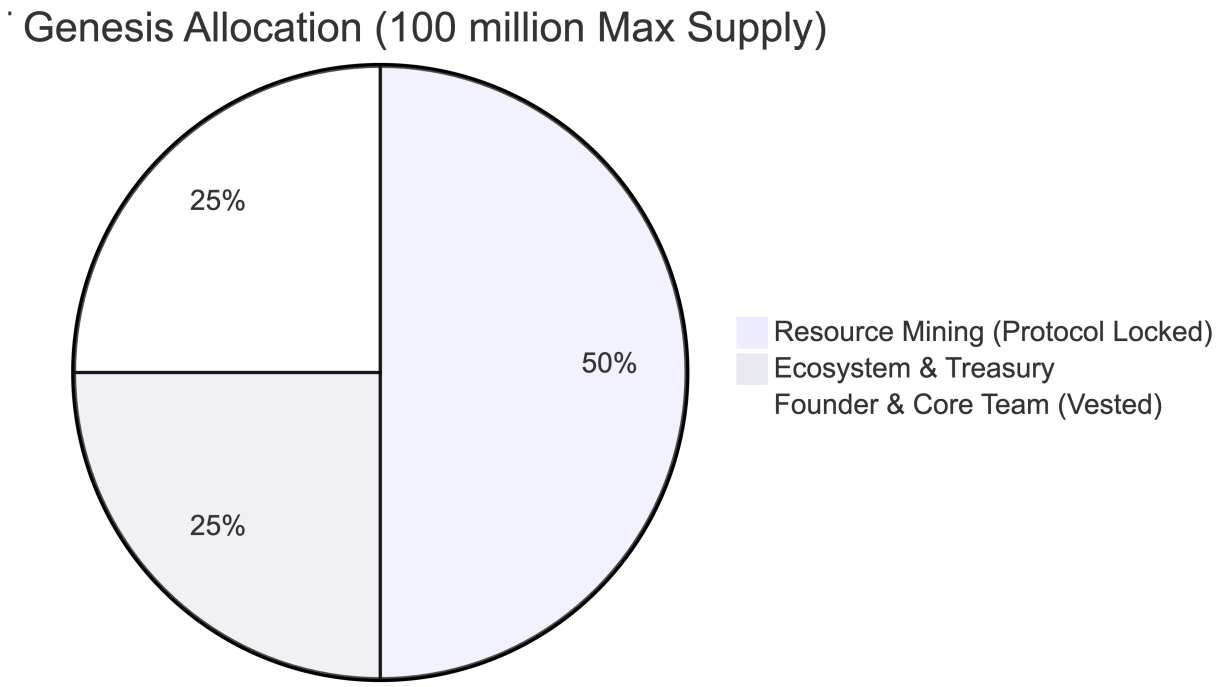

11. Applications of the Sovereign Fluid Grid

Just as Ethereum generalized state transitions to enable DeFi, TET Network generalizes thermodynamic compute to enable a new paradigm of AI-native applications.

- **Decentralized AI Inference Markets:** Startups can route LLM queries through TET at a fraction of centralized cloud costs, with execution mathematically verified by the ZK-Court.
- **Autonomous AI Agents (DAOs 2.0):** Smart contracts can natively invoke PoC nodes to run localized AI models, analyze markets, and execute trading or governance decisions independently without human intervention.
- **Quantum-Secure Financial Primitives:** Traditional DeFi applications built on TET inherit absolute post-quantum security, positioning the network as the primary settlement layer for high-value financial institutions in the quantum era.

Neural State Transition (The World Brain): Traditional blockchains store static balances; TET Network stores intelligence. During idle periods, PoC nodes utilize their spare compute to continuously fine-tune a globally shared, open-source base model via Continuous Federated Learning. The cryptographic state of the chain itself becomes a living, growing neural network. It is immune to corporate censorship and constantly evolves using privacy-preserving, real-world telemetry from up to 8 billion edge nodes.

Sentient Assets (Smart Contracts 2.0): By integrating the Neural State Transition directly into the runtime, TET enables the creation of digital assets with autonomous intelligence. Instead of rigid "if-then" logic, developers can deploy Sentient Assets—such as sovereign wallets that dynamically analyze transaction context to block sophisticated social engineering attacks, or digital entities that autonomously negotiate their own market value based on real-time ecosystem sentiment.

12. Development Roadmap

The deployment of the TET Network follows a phased architectural release:

- **Phase 0 (The Inference Wedge):** Launch of a lean, highly optimized LLM inference market targeting edge nodes. The protocol establishes the baseline energy-to-compute ratio using a standardized Llama-3 model.
- **Phase 1 (CAAC Implementation):** Full deployment of Context-Aware Adaptive Consensus. The network begins autonomously categorizing nodes into PoC and PoR roles based on real-time telemetry. ZK-Court dispute resolution is activated.
- **Phase 2 (Post-Quantum Fluid Grid):** Complete integration of ML-DSA signatures. The network achieves absolute fluid distribution, dynamically routing AI requests across a global, quantum-resistant matrix of up to 8 billion devices.

13. Threat Models and Cryptographic Defenses

To establish a truly sovereign L1, the protocol must be mathematically resilient against Byzantine actors and rational economic attacks.

- **Lazy Evaluation (Compute Spoofing):** A malicious node may attempt to save electricity by skipping the AI inference and submitting a fabricated result. TET Network defends against this via the ZK-Court. During the challenge window, Watchers can trigger a RISC Zero/SP1 zkVM execution trace. If the trace proves the initial commit was invalid, the malicious node suffers a 100% slashing of their staked $TET.

- **Hardware Spoofing (Sybil Attack):** A node might attempt to emulate a high-end GPU (e.g., RTX 4090) using lower-tier hardware to claim higher PoC rewards. TET mitigates this through *Probabilistic Hardware Fingerprinting*. The protocol issues non-deterministic micro-tasks that exploit specific hardware execution quirks, creating an unforgeable cryptographic proof of the physical hardware layer.

13.1 The Mathematics of Economic Finality

To deter Byzantine actors, TET Network enforces a Cost of Corruption (CoC)model that strictly exceeds the Profit from Attack(PfA). The slashed collateral S is defined as a multiplier λ(Lambda) of the expected reward $R_{expected}$:

$$S = \lambda \cdot R_{expected}$$

Because$\lambda \gg 1$, the economic expected value of any spoofing attempt is permanently locked in a negative state $(CoC > PfA)$. We do not trust the node; we trust the math that makes fraud irrational.

14. Conclusion

TET Network resolves the structural contradictions of contemporary Layer 1 protocols. By replacing homogeneous competition with Context-Aware Adaptive Consensus, the network eliminates hardware monopolies and maximizes thermodynamic efficiency. The integration of post-quantum cryptography ensures the mathematical survival of the ledger against future computational threats. TET Network is not merely a decentralized ledger; it is a fluid, autonomous infrastructure that redefines artificial intelligence compute and physical energy as a secure, sovereign, and universally accessible resource.

15. Join the Infrastructure (Call to Builders)

I have laid out the definitive blueprint for TET Network. The transition from theory to a sovereign infrastructure requires "Builders"—the 0.1% of engineers and cryptographers who possess the technical mastery to translate these formulas into a resilient codebase.

We are seeking Core Contributors proficient in low-level systems (**Rust**), decentralized networking (**libp2p**), and Applied Cryptography (**zkVMs, ML-DSA**).

No Résumés: We do not want your work history. If you possess the expertise to challenge or improve this architecture, submit a Technical Critique of this paper—identifying inefficiencies or proposing superior mechanisms—to **yizhenxianshi@gmail.com** with the subject line `[Core Builder Application]`.

Help us build the first block of the global resource grid.

16. References

[Core Consensus & Cryptoeconomics]

[1] Nakamoto, S. (2008). *Bitcoin: A Peer-to-Peer Electronic Cash System*.

[2] Buterin, V. (2014). *Ethereum: A Next-Generation Smart Contract and Decentralized Application Platform*.

[3] Sompolinsky, Y., & Zohar, A. (2015). *Secure High-Rate Transaction Processing in Bitcoin*. (Financial Cryptography and Data Security).

[4] Buterin, V., & Griffith, V. (2017). *Casper the Friendly Finality Gadget*. (Ethereum Foundation - Cost of Corruption and Economic Finality).

[Zero-Knowledge & Post-Quantum Cryptography]

[5] Ben-Sasson, E., Chiesa, A., et al. (2014). *SNARKs for C: Verifying Program Executions Succinctly and in Zero Knowledge*. (CRYPTO 2013).

[6] RISC Zero Team. (2023). *RISC Zero: A Zero-Knowledge Virtual Machine for General Purpose Computing*.

[7] Succinct Labs. (2024). *SP1: A Next-Generation zkVM for High-Performance Proof Generation*. [8] National Institute of Standards and Technology (NIST). (2024). *FIPS 204: Module-Lattice-Based Digital Signature Standard (ML-DSA)*.

[Decentralized AI & Network Layer]

[9] Borzunov, A., et al. (2022). *Petals: Collaborative Inference and Fine-tuning of Large Models*. (arXiv:2209.01188).

[10] Rao, Y. (2021). *Bittensor: A Peer-to-Peer Intelligence Market*.

[11] AI@Meta. (2024). *Llama 3 Model Card*. (Meta AI Research).

[12] Protocol Labs. (2019). *libp2p: A Modular Network Stack for Peer-to-Peer Applications*.

[Thermodynamics & Hardware Security]

[13] McMahan, B., et al. (2017). *Communication-Efficient Learning of Deep Networks from Decentralized Data*. (AISTATS 2017 - The Foundation of Federated Learning).
[14] Landauer, R. (1961). *Irreversibility and Heat Generation in the Computing Process*. (IBM Journal of Research and Development).

[15] Bennett, C. H. (1982). *The Thermodynamics of Computation – Individual and Society*.

[16] Suh, G. E., & Devadas, S. (2007). *Physical Unclonable Functions for Device Authentication and Secret Key Generation*. (DAC 2007).